Introduction to Hadoop

Assoc. Prof. Dr. Thanachart Numnonda
Executive Director
IMC Institute
April 2016

What is Hadoop?

A scalable fault-tolerant distributed system for data storage and processing

Completely written in java
Open source & distributed under Apache license

Hadoop Environment

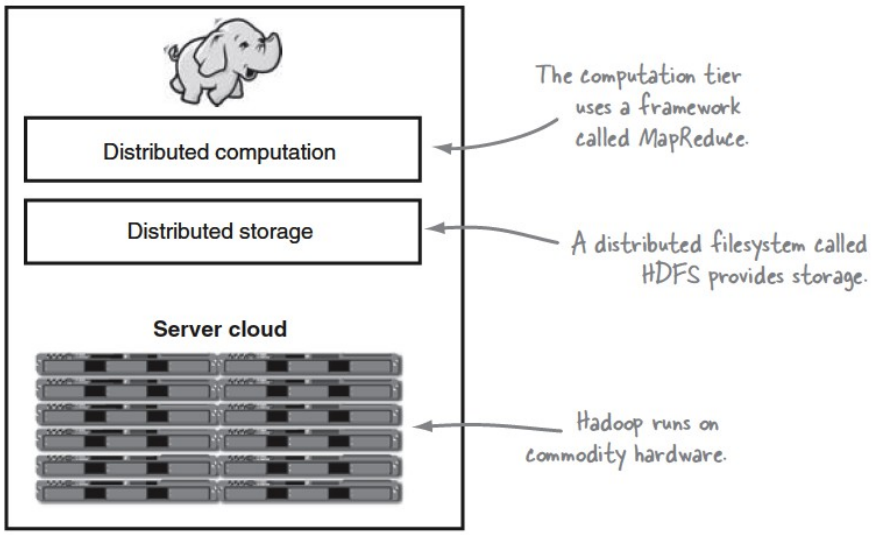

Source: Hadoop in Practice; Alex Holmes

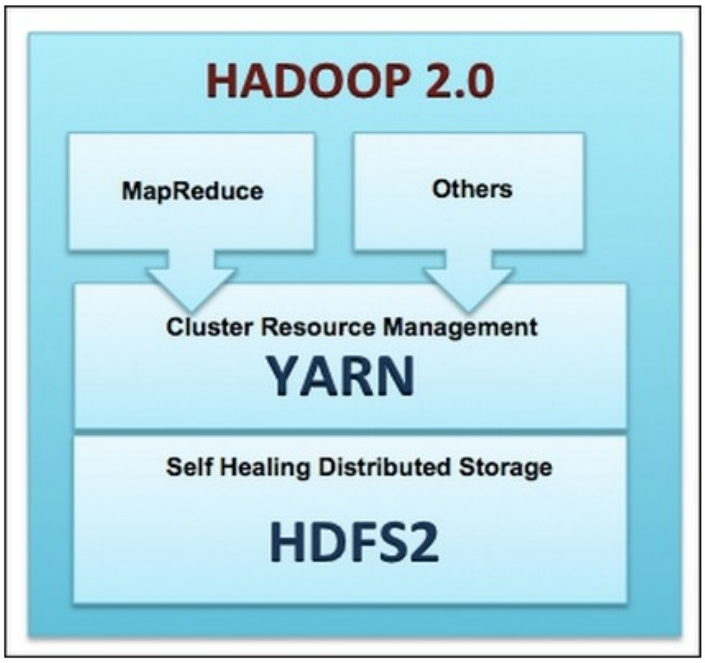

Source: HDInsight Essentials - Second Edition

Hadoop Ecosystem

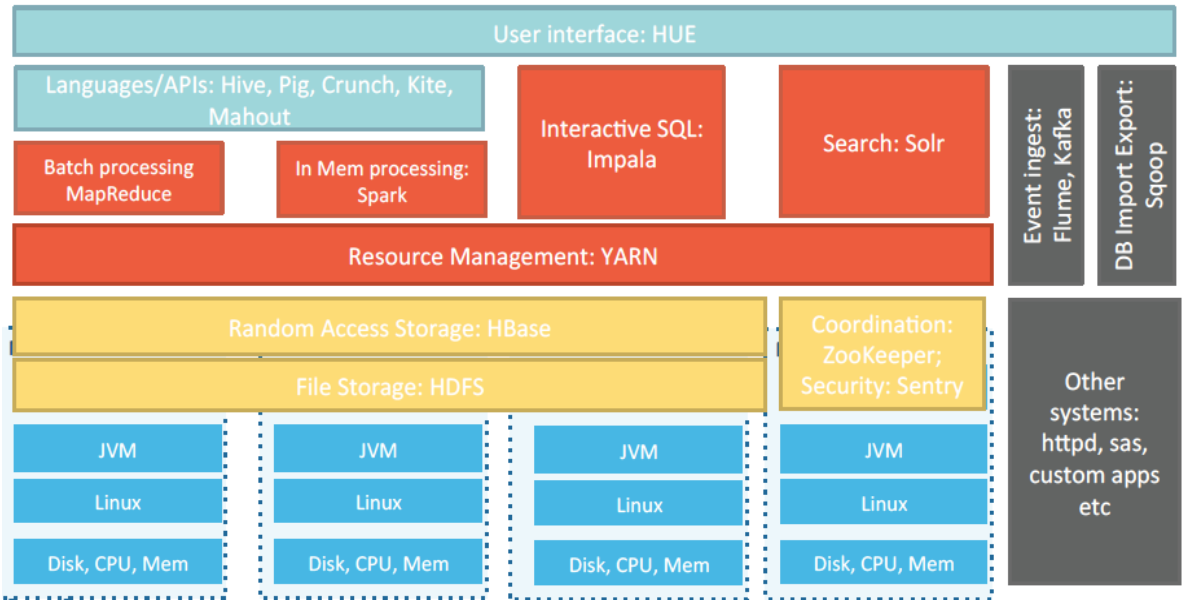

Source: Apache Hadoop Operations for Production Systems, Cloudera

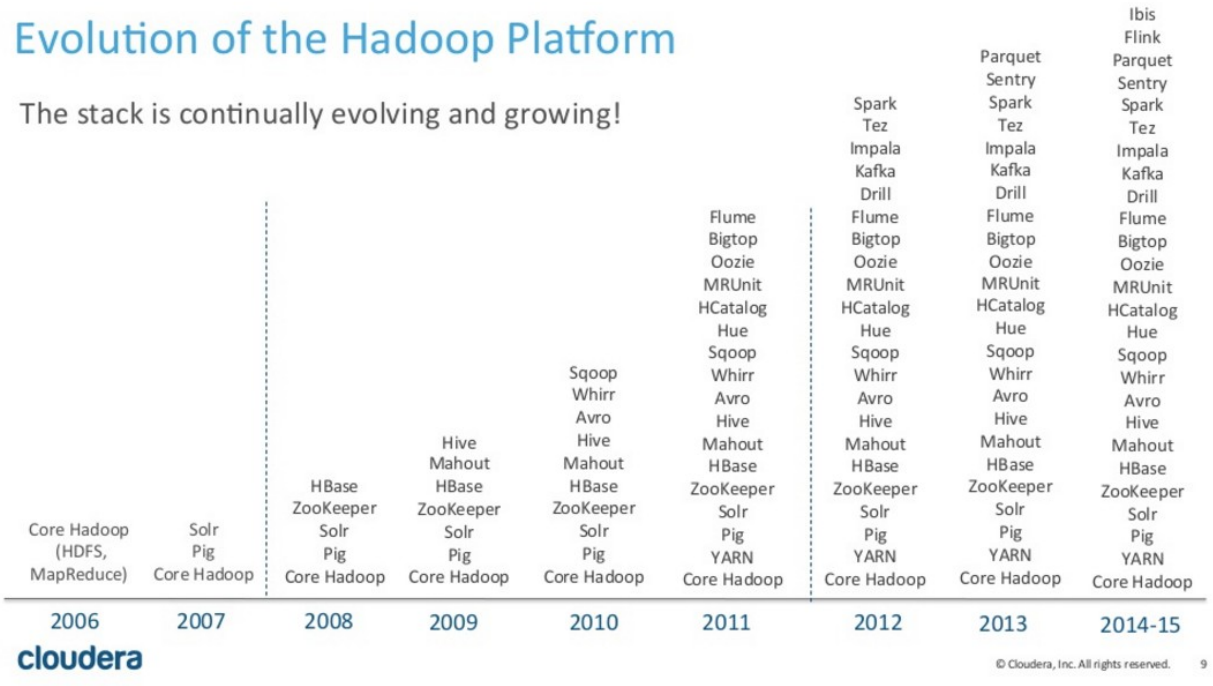

Source: The evolution and future of Hadoop storage: Cloudera

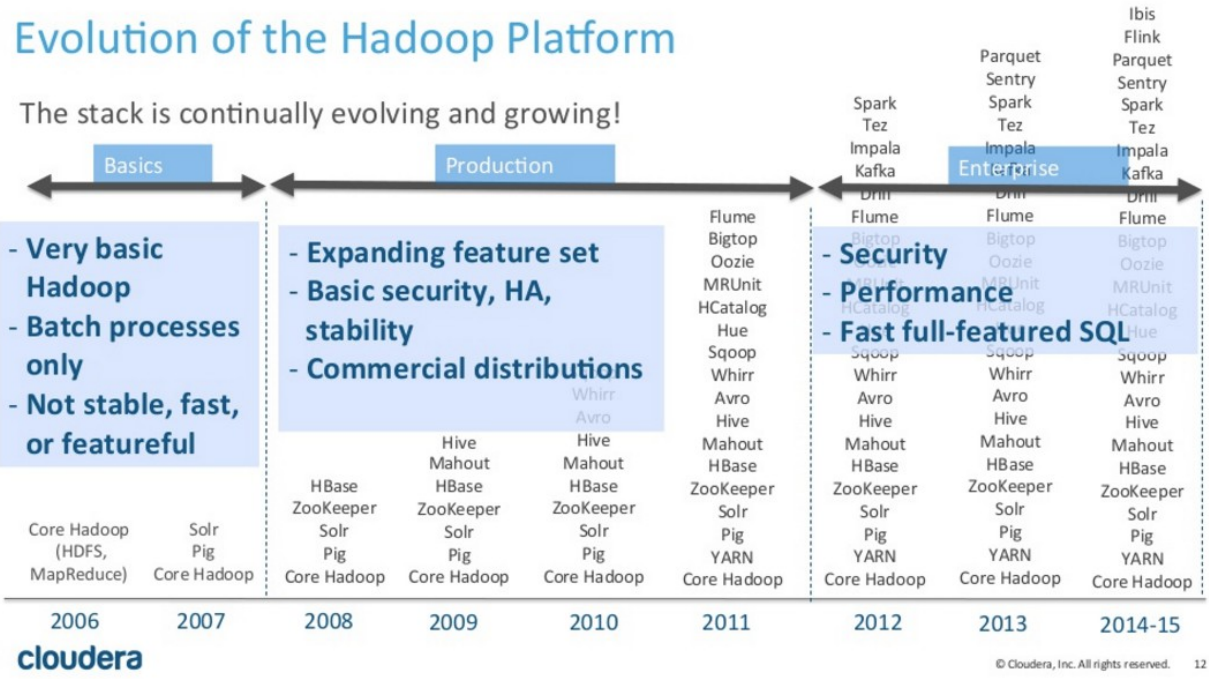

Source: The evolution and future of Hadoop storage: Cloudera

Hadoop Cluster

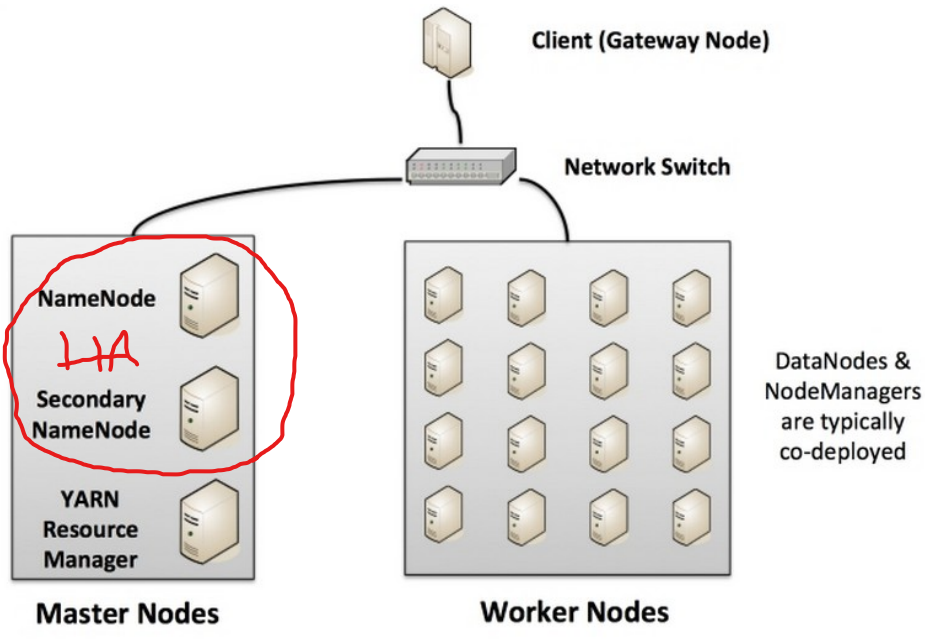

Source: HDInsight Essentials - Second Edition

Data Processing Technology using Hadoop

- MapReduce
- Hive
- Pig
- Imapla
- Drill
- Spark

What is MapReduce?

MapReduce is a programming model and an associated implementation for processing and generating large data sets with a parallel, distributed algorithm on a cluster

A MapReduce program is composed of a Map() method and a Reduce() method. It runs on a batch mode.

What is Hive?

Hive is developed by Facebook,
designed to enable easy data summarization,
ad-hoc querying and analysis of large volumes of data.

It provides a simple query language called Hive QL,
which is based on SQL.
It is not designed for online transaction processing
and does not offer real-time queries
and row level updates

What is Pig?

Pig is a platform for analyzing large data sets
that consists of a high-level language
for expressing data analysis programs,
coupled with infrastructure for evaluating these programs.

It is developed by Yahoo and uses a scripting language
Called Pig Latin.

What is Impala?

Impala is a query engine that runs on Apache Hadoop.
It brings scalable parallel database technology
to Hadoop, enabling users to issue
low-latency SQL queries to data stored in HDFS
and Apache HBase without requiring
data movement or transformation.

What is Spark?

An open source big data processing framework
built around speed, ease of use,
and sophisticated analytics.

Spark enables applications in Hadoop clusters
to run up to 100 times faster in memory
and 10 times faster even when running on disk.

Open Source Ecosystem

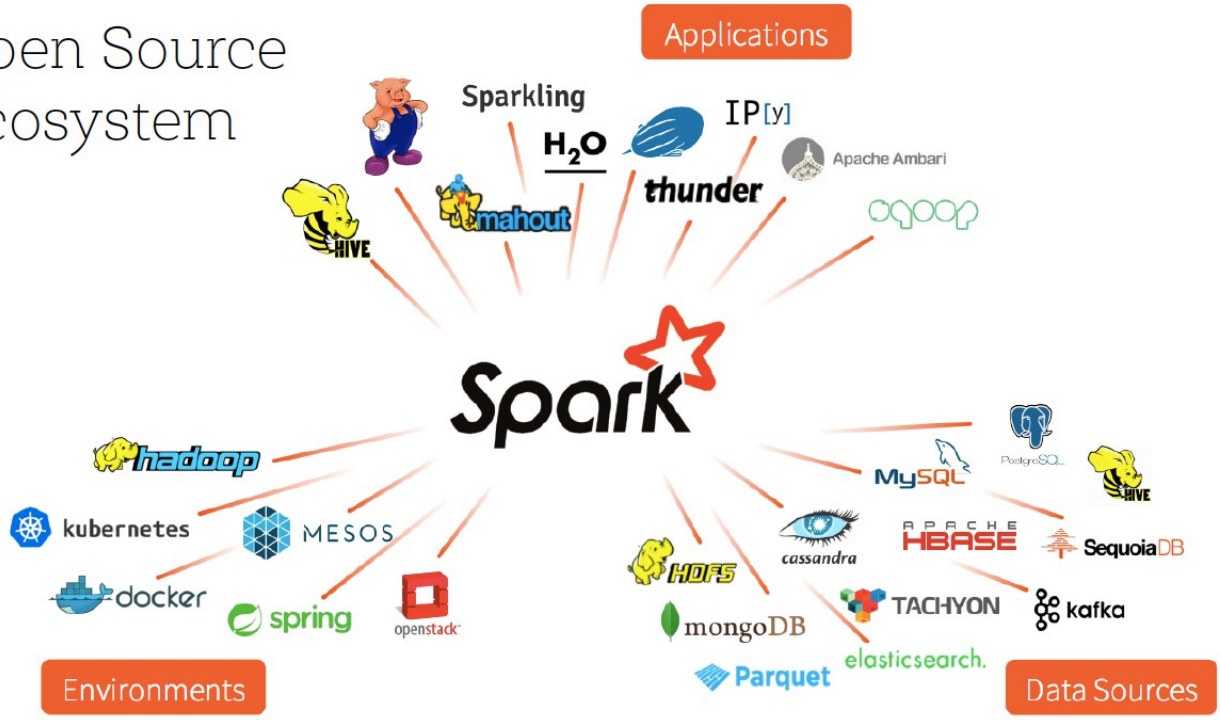

Source: Jump start into Apache Spark and Databricks

Processing Techniques

	Hive	Pig	Spark	Impala
Query Latency	Medium (Tez), High	High	Low	Low
Storage	HDFS	HDFS	HDFS,Object Storage	HDFS
Language	Hive QL	Pig Latin	Java, Python, Scala, R	Impala
Analysis	Batch	Batch	Interactive, Real-time (Spark Streaming), Machine Learning (MLlib)	Interactive
SQL Compatibility	Medium	No	Low (Spark SQL)	Medium

Lisp
Haskell

What is HDFS?

**A distributed file system that provides high-performance access to data across Hadoop clusters.
Like other Hadoop-related technologies, HDFS has become a key tool for managing pools of big data and supporting big data analytics applications.**

HDFS Architecture

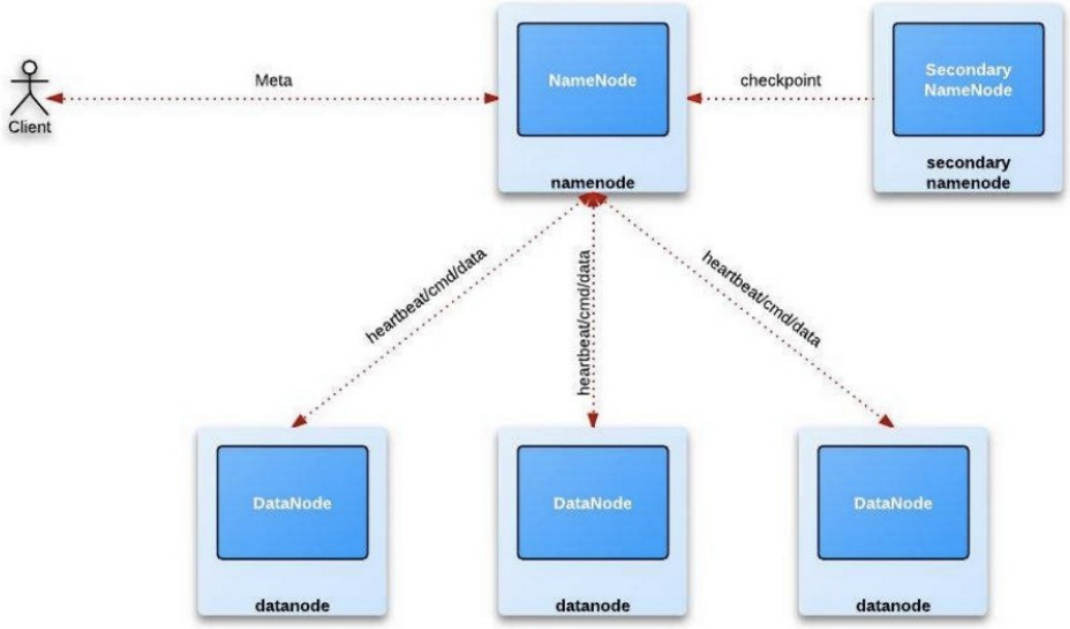

Source: Details intro to HDFS, Clay Jiang

HDFS Storage

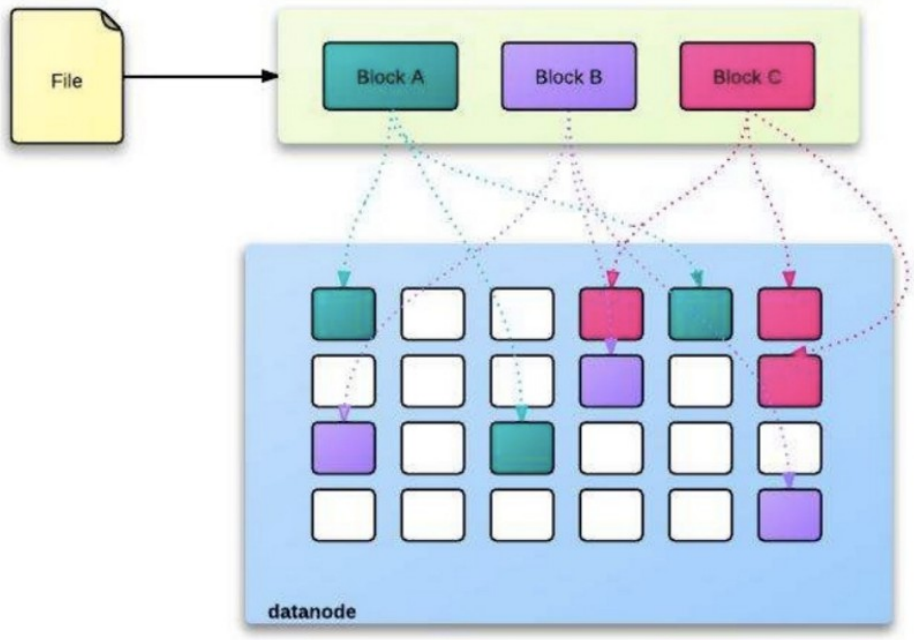

Source: Details intro to HDFS, Clay Jiang

What is HBase?

HBase is an open source, non-relational, distributed database modeled after Google's BigTable and is written in Java. It is developed as part of Apache Software Foundation's Apache Hadoop project and runs on top of HDFS (providing BigTable-like capabilities for Hadoop). That is, it provides a fault-tolerant way of storing large quantities of sparse data.

APACHE HBASE

HBase Components

- Region
  - Row of table are stores
- Region Server
  - Hosts the tables
- Master
  - Coordinating the Region Servers
- ZooKeeper
- HDFS
- API
  - The Java Client API

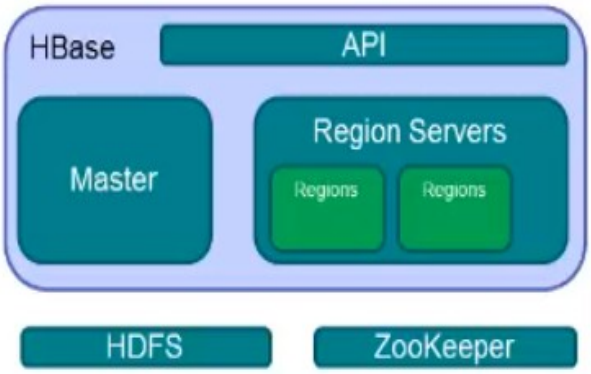

Hive.apache.org

Evolution of Storage (Basics / 2006-2007)

- **HDFS only**
- Support basic batch workloads. No HA.
- Performance not important
  - MapReduce is too slow, anyway!
  - Batch only
- Early Adopters (FaceBook, Yahoo, etc)

Source: The evolution and future of Hadoop storage: Cloudera

Evolution of Storage (Production / 2008-2011)

- **HDFS evolves** to add high availability and security
  - Focused on batch workloads
  - Inefficient file formats commonly used (text)
  - Query engines are slow! No need for better performance
- **Apache HBase** becomes an Apache Top-Level Project (TLP)
  - **Introduces fast random access**
  - Early adopters experiment with new use cases
  - Deployed at Facebook and other large companies

Source: The evolution and future of Hadoop storage: Cloudera

Evolution of Storage (Production / 2008-2011)

- **HDFS evolves** to add high availability and security
  - Focused on batch workloads
  - Inefficient file formats commonly used (text)
  - Query engines are slow! No need for better performance
- **Apache HBase** becomes an Apache Top-Level Project (TLP)
  - **Introduces fast random access**
  - Early adopters experiment with new use cases
  - Deployed at Facebook and other large companies

Source: The evolution and future of Hadoop storage: Cloudera

Evolution of Storage (Enterprise / 2012-2015)

- **Reliable core** brings new users
  - **Enterprise features**: access control, disaster recovery, encryption
- Introduction of **fast query engines**
  - 10-100x faster **SQL-on-Hadoop** (Impala, Spark, etc.)
  - Pushes **HDFS performance improvements**: caching, CPU efficiency, columnar file formats (Apache Parquet, ORCFile)
- **HBase evolves to 1.0**
  - Improved stability, scalability, security
  - Good **random** access - not fast for SQL analytics.
- **Initial support for cloud storage**
  - Rising adoption of AWS, Azure, Google Compute, etc.

Source: The evolution and future of Hadoop storage: Cloudera

2016-2020 (Next-gen): storage hardware

- **Spinning disk -> solid state storage**
  - **NAND flash**: Up to 450k read 250k write iops, about 2GB/sec read and 1.5GB/sec write throughput, at a price of less than $3/GB and dropping fast
  - **3D XPoint memory** (1000x faster than NAND, cheaper than RAM)
- **RAM** is cheaper and more abundant:
  - 64->128->256GB over last few years
- **HDFS and HBase** were not designed for next-generation hardware.
  - Not using full speed of flash or RAM size

Source: The evolution and future of Hadoop storage: Cloudera

2016-2020 (Next-gen): gaps in capabilities

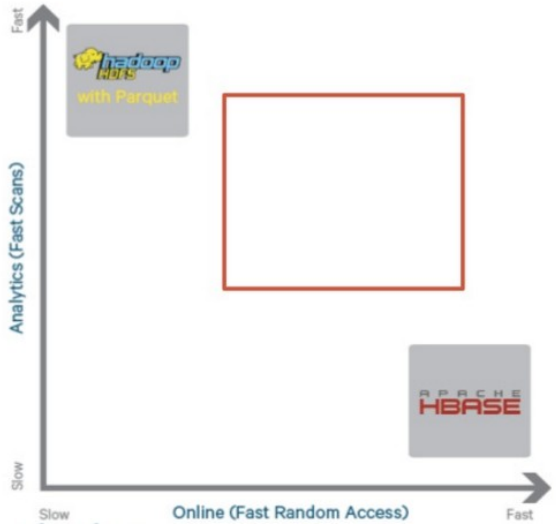

HDFS good at:

- Batch ingest only (eg hourly)
- Efficiently scanning large amounts of data (analytics)

HBase good at:

- Efficiently finding and writing individual rows
- Making data mutable

Gaps exist when these properties are needed *simultaneously*

Source: The evolution and future of Hadoop storage: Cloudera

2016-2020 (Next-gen): Apache Kudu (incubating)

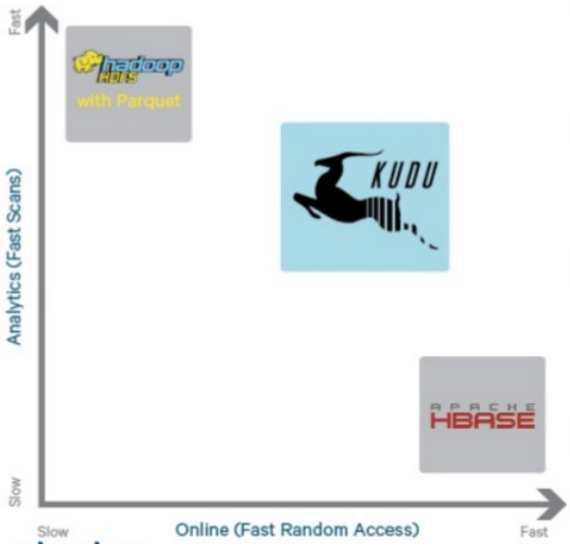

- **High throughput** for big scans
  Goal: Within 2x of Parquet
- **Low-latency** for short accesses
  Goal: 1ms read/write on SSD
- **Relational data model**
  - SQL queries are easy
  - "NoSQL" style scan/insert/update (Java/C++ client)
- **Expands Hadoop use cases**
  - Real-time analytics and time series
  - Internet-of-things

Source: The evolution and future of Hadoop storage: Cloudera

Kudu: Open source, scalable and fast tabular storage

- **Scalable**
  - Designed to scale to **1000s of nodes, tens of PBs**
- **Fast**
  - **Designed for modern hardware**
  - **Millions** of read/write operations per second across cluster
  - **Multiple GB/second** read throughput per node
- **Tabular**
  - Store tables like a normal database (**support SQL, Spark, etc**)
  - NoSQL-style access to **100+ billion row tables** (Java/C++/Python APIs)

Source: The evolution and future of Hadoop storage: Cloudera

Kudu Use Cases

Kudu is best for use cases requiring a simultaneous combination of sequential and random reads and writes

- **Time Series**
  - Examples: Stream market data; fraud detection & prevention; risk monitoring
  - Workload: Insert, updates, scans, lookups
- **Online Reporting**
  - Examples: ODS
  - Workload: Inserts, updates, scans, lookups

Source: Apache Kadu: Fast Analytics and Fast Data: Cloudera

Real-Time Analytics in Hadoop Today

Fraud Detection in the Real World = Storage Complexity

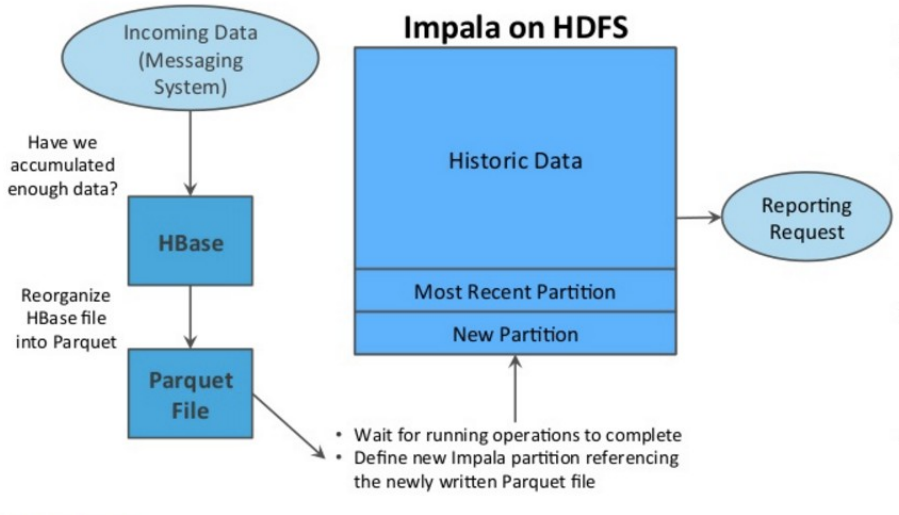

Considerations:

- How do I handle failure during this process?
- How often do I reorganize data streaming in into a format appropriate for reporting?
- When reporting, how do I see data that has not yet been reorganized?
- How do I ensure that important jobs aren't interrupted by maintenance?

Source: Apache Kadu: Fast Analytics and Fast Data: Cloudera

2016-2020 (Next gen): Predictions

- **Kudu** will evolve an enterprise feature set and enable simple high-performance real-time architectures
  - Increasing ability to migrate traditional applications
- **HDFS and HBase** will continue to innovate and adapt to next generation hardware
  - Steady improvements in **performance, efficiency, and scalability** (e.g. erasure coding)
- **Cloud storage** will become increasingly important
  - **Hadoop ecosystem** will evolve to coexist

Source: The evolution and future of Hadoop storage: Cloudera

Other Hadoop Services

- **Flume:** Collection and import of log and event data
- **Kafka**: Publish-subscribe messaging rethought as a distributed commit log. Fast.
- **Sqoop**: Imports data from relational databases
- **Oozie**: Hadoop workflow.
- **Mahout**: Machine Learning Libray
- **Zookeeper**: Centralized service for maintaining configuration
- **Ambari**: Cluster management

Hadoop Ecosystem

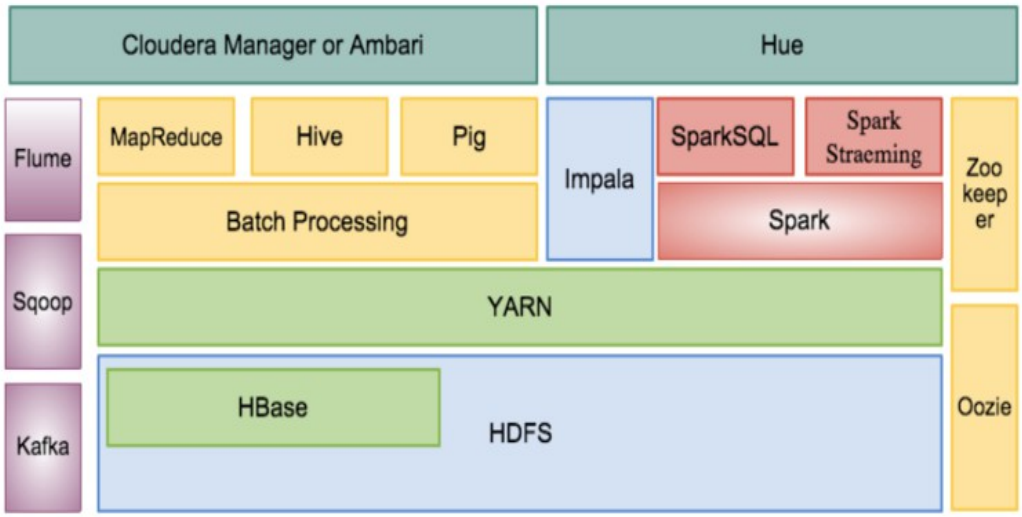

Use Case

Lamda Arch.

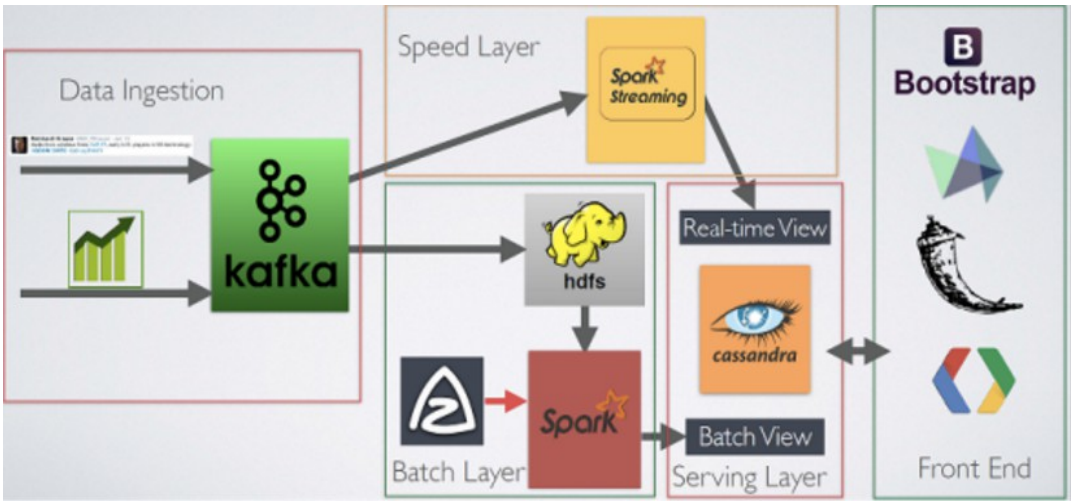

Source: http://www.insightdataengineering.com/

Hadoop Cluster

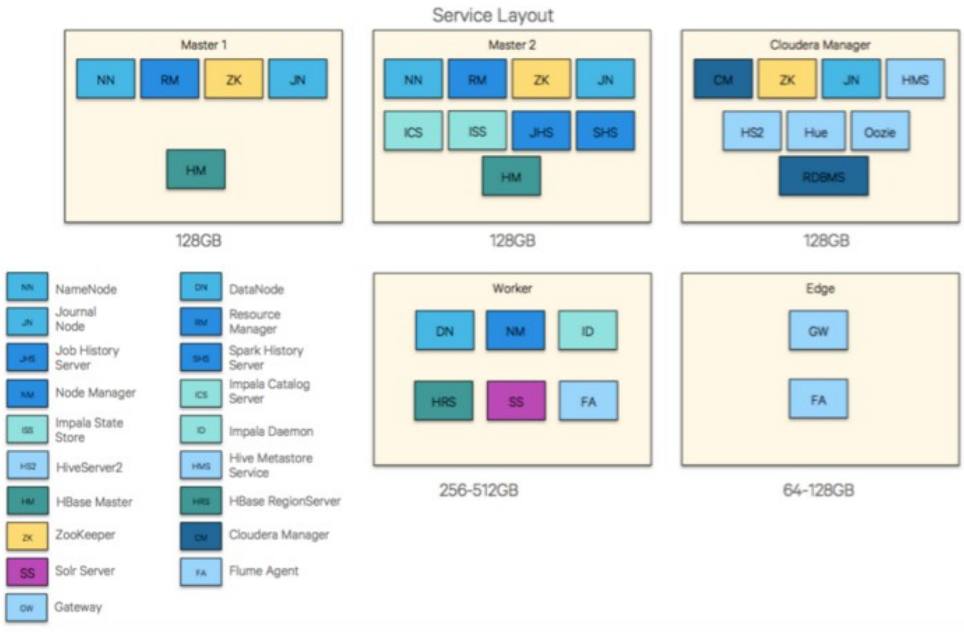

Source: How-to: Deploy Apache Hadoop Clusters Like a Boss, Jeff Holoman

Hadoop Hardware Sizing

CPU	2 × 6 core 2.9 Ghz/15 MB cache
Memory	64 GB DDR3-1600 ECC
Disk controller	SAS 6 Gb/s
Disks	12 × 3 TB LFF SATA II 7200 RPM
Network controller	2 × 1 Gb Ethernet

Hadoop Cluster Size: Survey

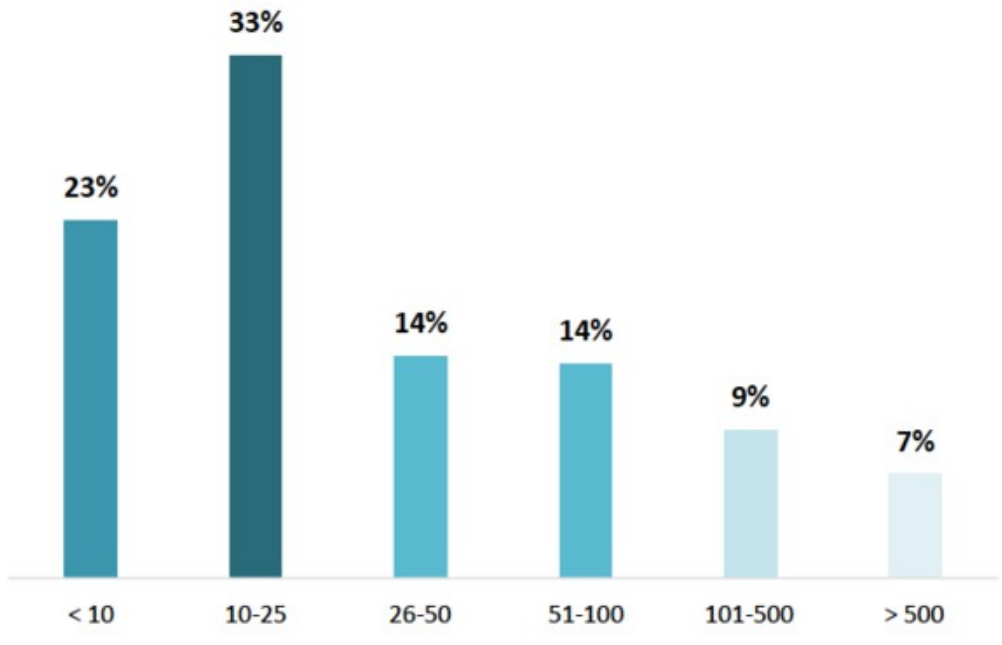

Source: AtScale

Hadoop Distribution

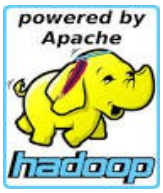

cloudera

Pivotal

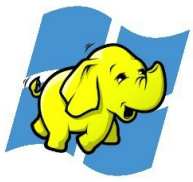

Azure HDInsight

The Forrester Wave™: Big Data Hadoop Distributions: Q1 2006

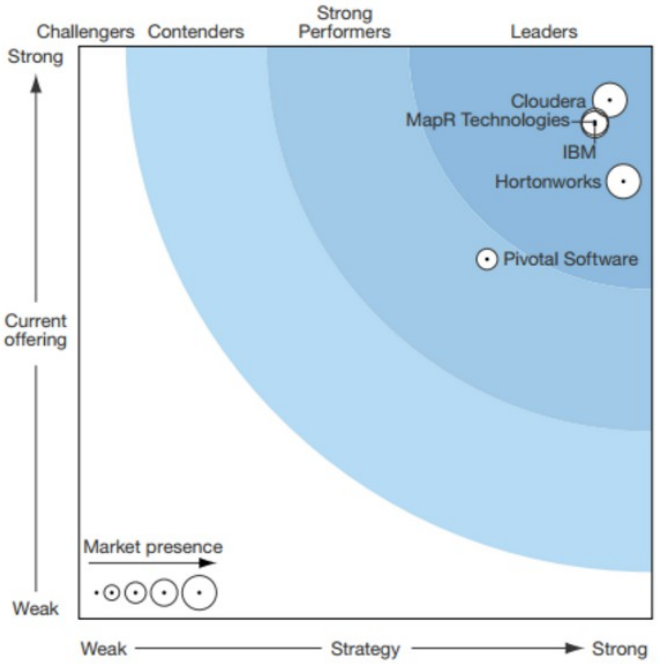

Cloudera

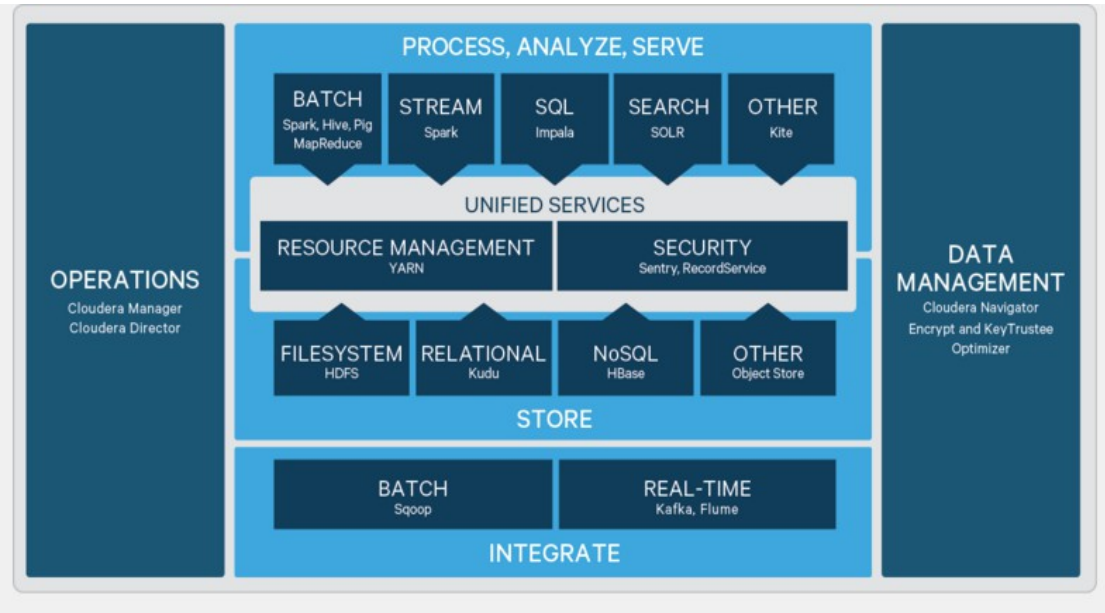

MapR

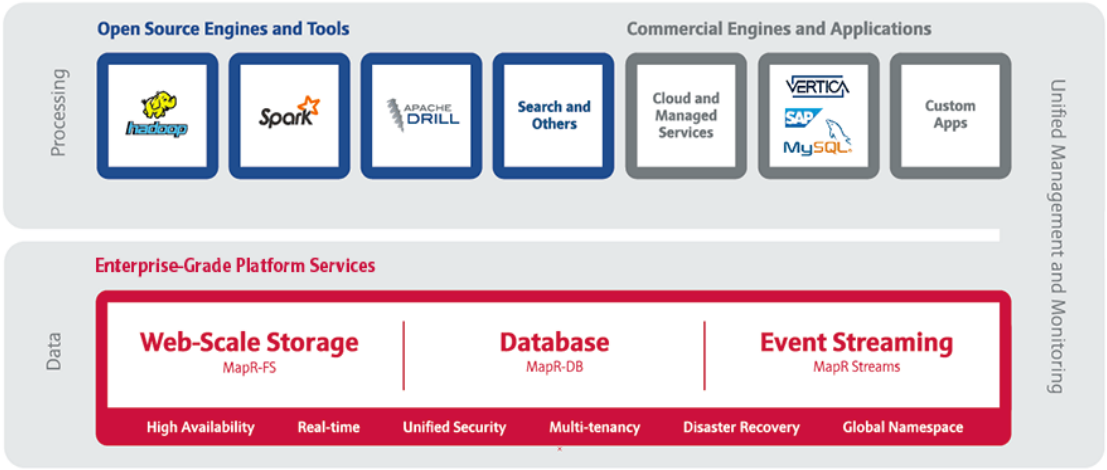

Hortonworks

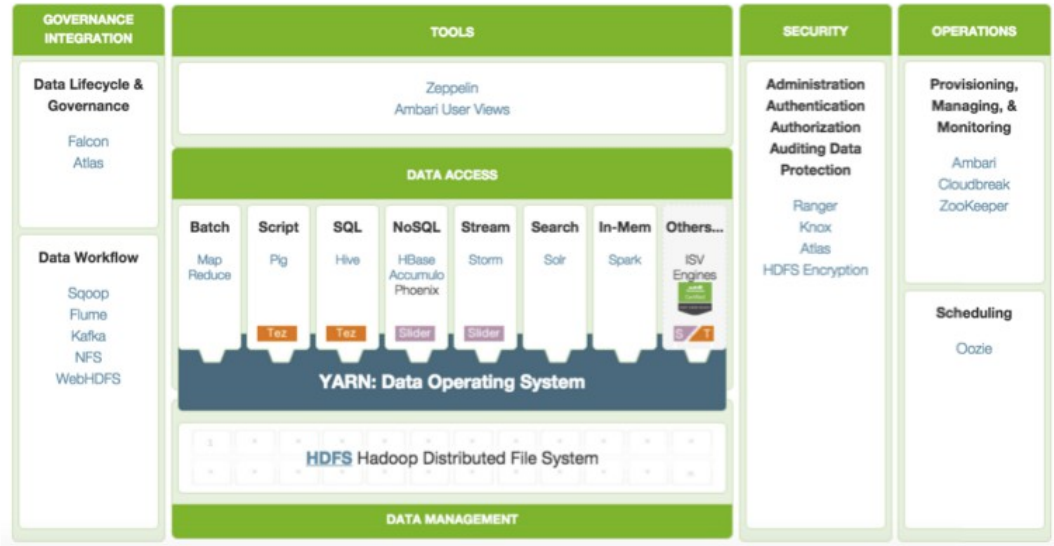

Hortonworks

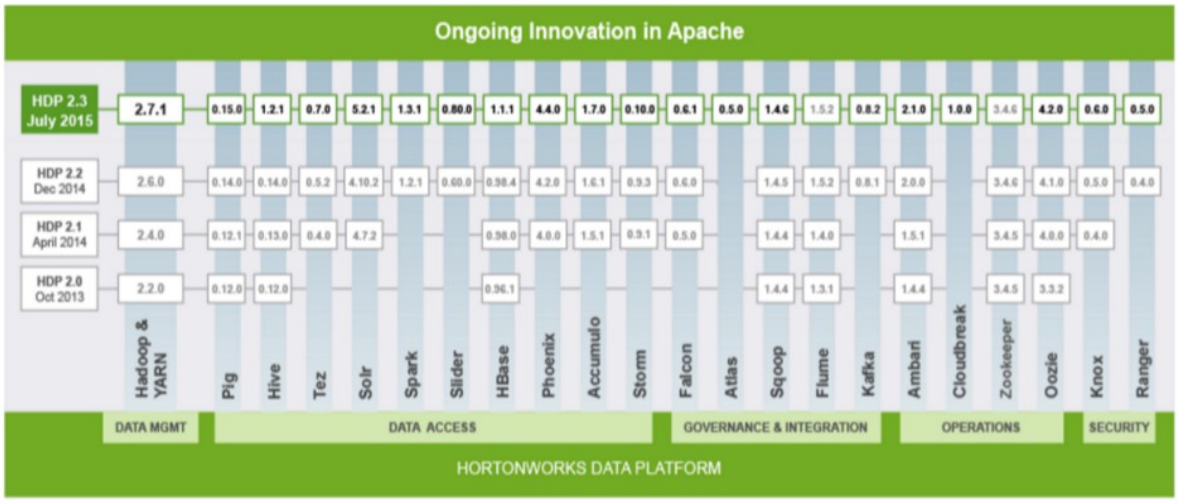

Pivotal HDB

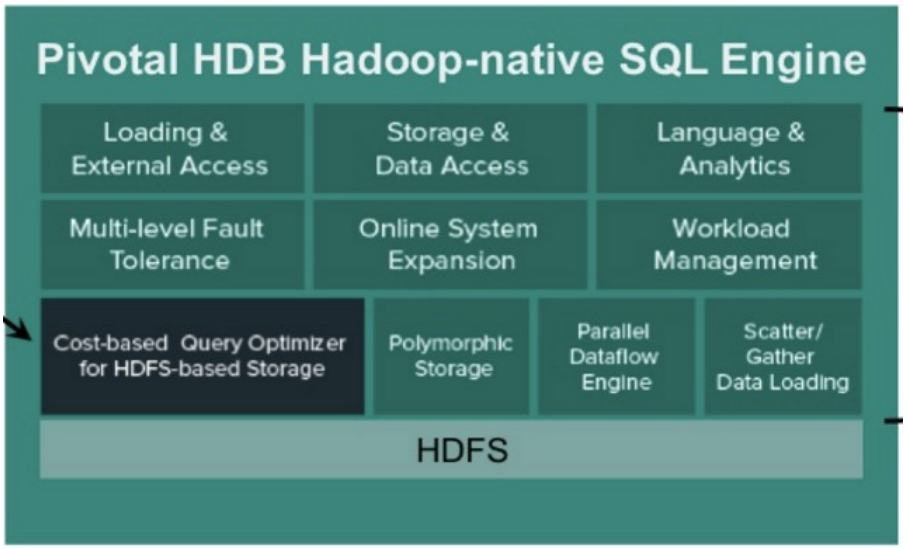
Pivotal HDB Hadoop-native SQL Engine
Loading & External Access
Storage & Data Access
Language & Analytics
Multi-level Fault Tolerance
Online System Expansion
Workload Management
Cost-based Query Optimizer for HDFS-based Storage
Polymorphic Storage
Parallel Dataflow Engine
Scatter/ Gather Data Loading
HDFS

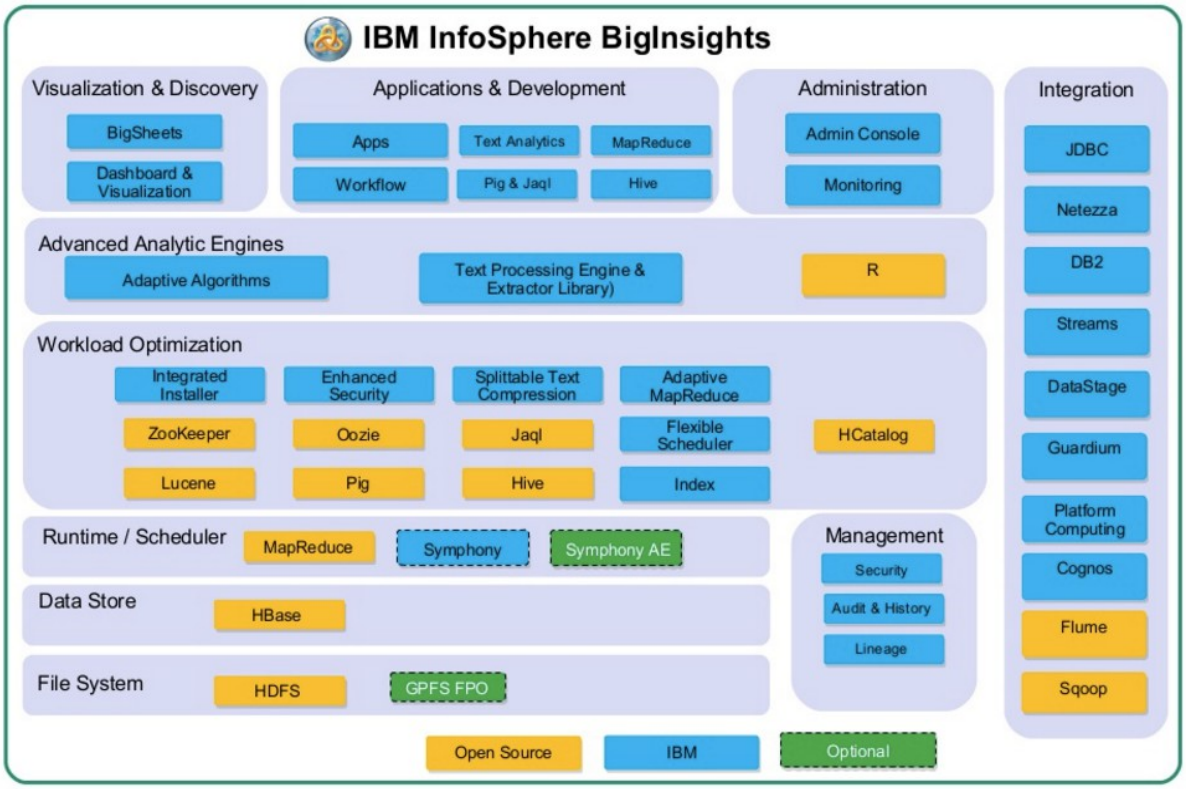
IBM InfoSphere BigInsights
Visualization & Discovery
BigSheets
Dashboard & Visualization
Applications & Development
Apps
Text Analytics
MapReduce
Workflow
Pig & Jaql
Hive
Administration
Admin Console
Monitoring
Integration
JDBC
Netezza
DB2
Streams
DataStage
Guardium
Platform Computing
Cognos
Flume
Sqoop
Advanced Analytic Engines
Adaptive Algorithms
Text Processing Engine & Extractor Library)
R
Workload Optimization
Integrated Installer
Enhanced Security
Splittable Text Compression
Adaptive MapReduce
ZooKeeper
Oozie
Jaql
Flexible Scheduler
HCatalog
Lucene
Pig
Hive
Index
Runtime / Scheduler
MapReduce
Symphony
Symphony AE
Management
Security
Audit & History
Lineage
Data Store
HBase
File System
HDFS
GPFS FPO
Open Source
IBM
Optional

Amazon EMR

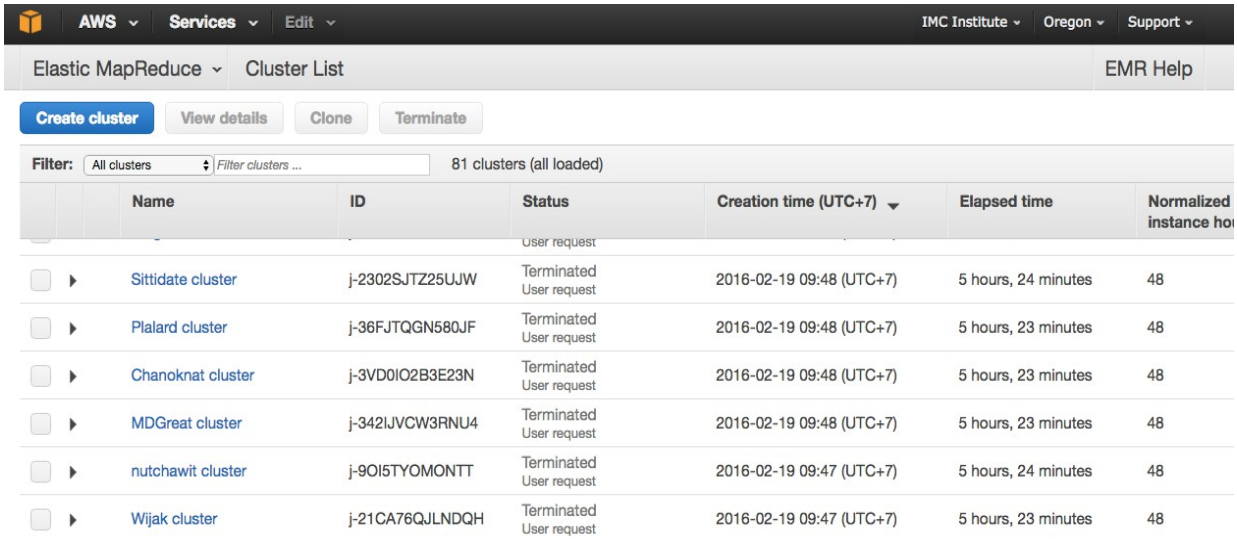

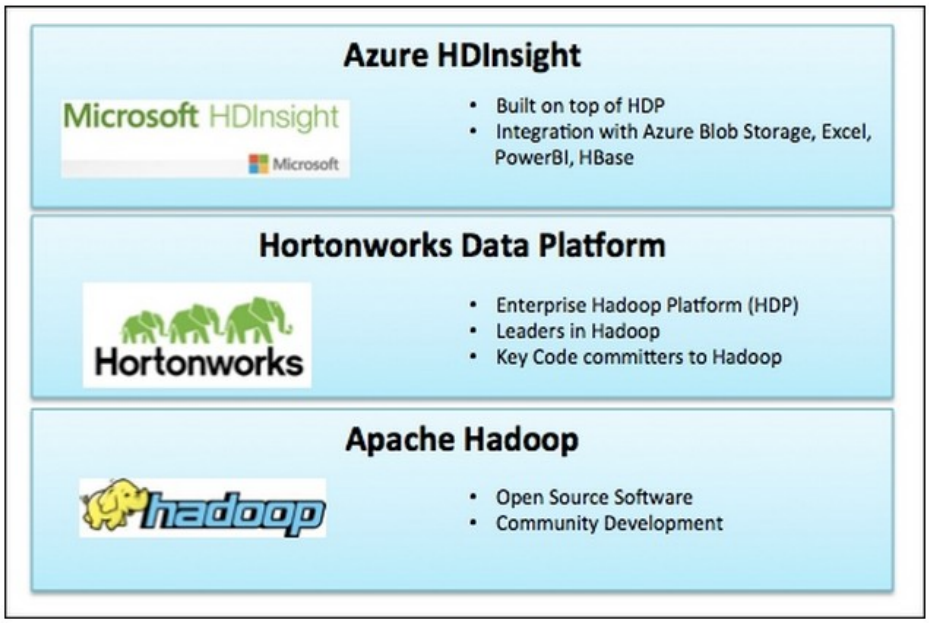

Source: HDInsight Essentials - Second Edition

www.facebook.com/imcinstitute

Thank you

thanachart@imcinstitute.com

www.imcinstitute.com

www.facebook.com/imcinstitute

www.slideshare.net/imcinstitute

www.thanachart.org